

Overview of neural network architectures for graph-structured data analysis

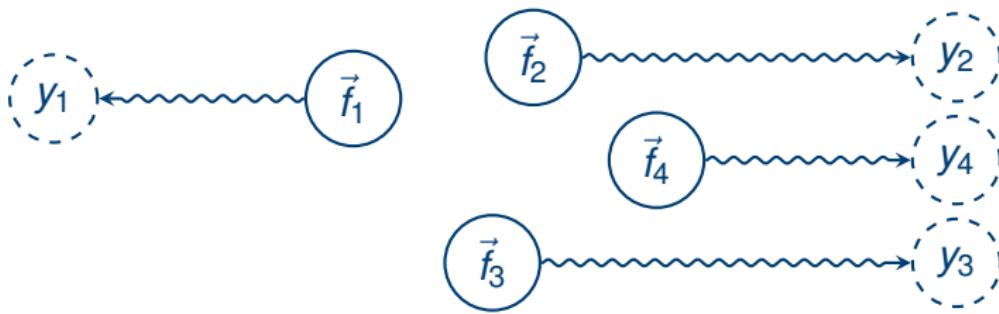
Petar Veličković

Artificial Intelligence Group
Department of Computer Science and Technology, University of Cambridge, UK

Motivation: supervised learning



- ▶ Petar Veličković here!
- ▶ This is a (supervised) machine learning problem.

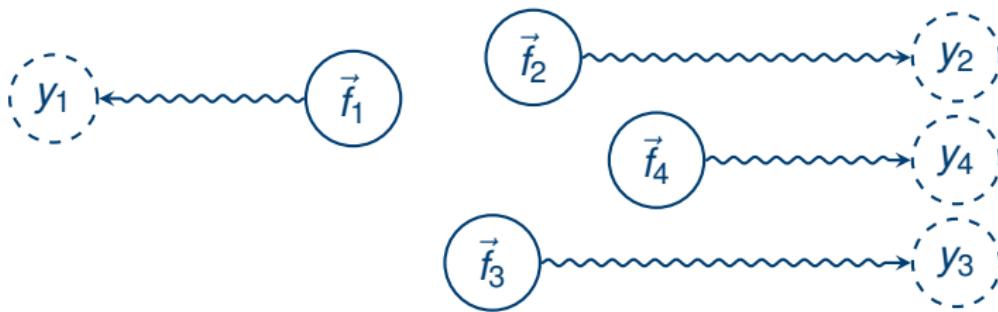


- ▶ Four examples, features (\vec{f}_i) and labels (y_i).
- ▶ Good enough for science. ✓

Motivation: supervised learning



- ▶ Petar Veličković here!
- ▶ This is a (supervised) machine learning problem.

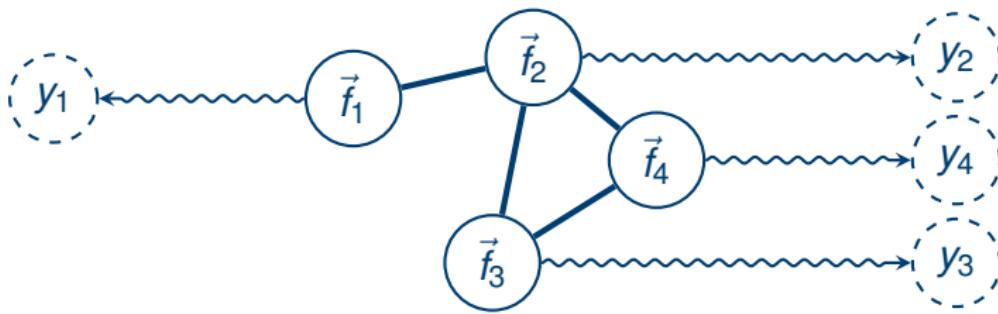


- ▶ Four examples, features (\vec{f}_i) and labels (y_i).
- ▶ Good enough for science. **Not Aperture Science!**.....X

Motivation: supervised learning

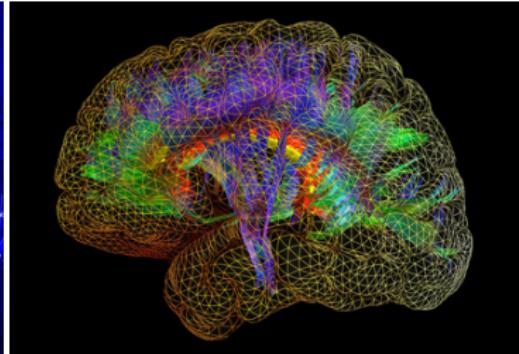
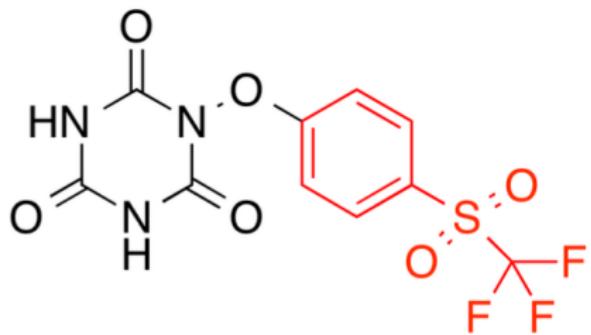


- ▶ Petar Veličković here!
- ▶ This is a (supervised) machine learning problem.



- ▶ Four examples, features (\vec{f}_i) and labels (y_i).
- ▶ Good enough for science. **Not Aperture Science!** X
- ▶ Gentlemen, I give you **graphs**. *The inputs of tomorrow!*

Graphs are **everywhere**!



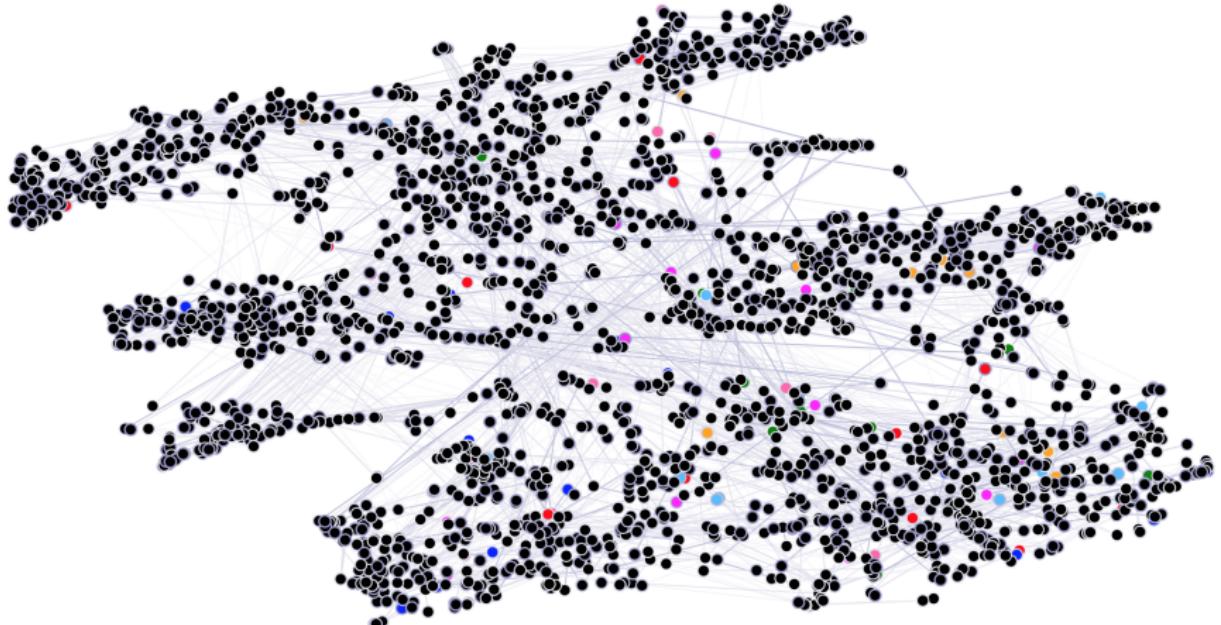
Introduction

- ▶ In this talk, I will demonstrate some of the popular methodologies that leverage **neural networks** for processing **graph-structured inputs**.
- ▶ Although the earliest approaches to this problem date to the late 90s, it has caught traction only in the recent five years (with a proper explosion happening throughout 2017)!
 - ▶ For early references, you may investigate the works of Sperduti & Starita (1997) and Frasconi *et al.* (1998), IEEE TNNLS.
- ▶ There's at least ten submissions *to ICLR 2018 alone* that attempt solving the same graph problems in different ways.

Mathematical formulation

- ▶ We will focus on the **node classification** problem:
 - ▶ **Input:** a matrix of *node features*, $\mathbf{F} \in \mathbb{R}^{N \times F}$, with F features in each of the N nodes, and an *adjacency matrix*, $\mathbf{A} \in \mathbb{R}^{N \times N}$.
 - ▶ **Output:** a matrix of *node class probabilities*, $\mathbf{Y} \in \mathbb{R}^{N \times C}$, such that $Y_{ij} = \mathbb{P}(\text{Node } i \in \text{Class } j)$.
- ▶ We also assume, for simplicity, that the edges are **unweighted** and **undirected**:
 - ▶ That is, $A_{ij} = A_{ji} = \begin{cases} 1 & i \leftrightarrow j \\ 0 & \text{otherwise} \end{cases}$
- ▶ There are **two** main kinds of learning tasks in this space...

Transductive learning



Training algorithm sees *all features (including test nodes)*!

Inductive learning

- ▶ Now, the algorithm *does not have access to all nodes upfront!*
- ▶ This often implies that either:
 - ▶ Test nodes are (incrementally) inserted into training graphs;
 - ▶ Test graphs are **disjoint** and *completely unseen!*
- ▶ A much harder learning problem (requires generalising across *arbitrary graph structures*), and many transductive methods will be inappropriate for inductive problems!

Simplest approach: a per-node classifier

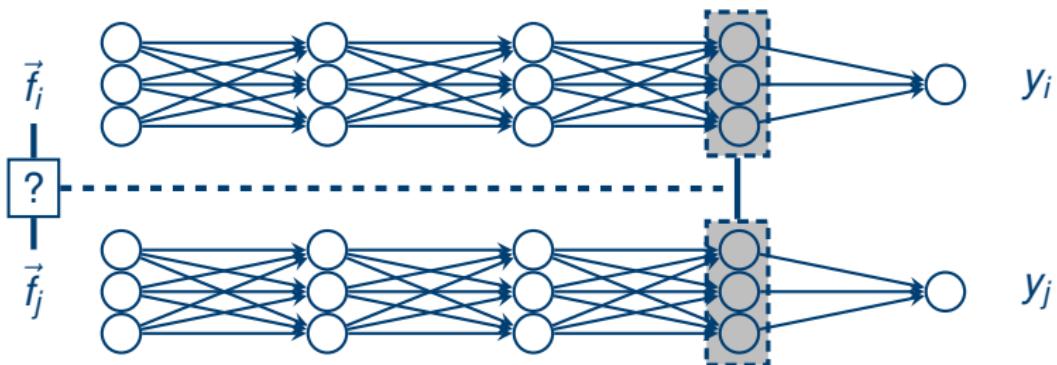
- ▶ Completely **drop** the graph structure, and classify each node individually, with a shared deep neural network classifier. :)
- ▶ In fact, this is how *most of deep learning is done*, even if there might be relationships between training examples!
- ▶ A single layer of the network computes $\mathbf{F}' = \sigma(\mathbf{FW})$, where $\mathbf{W} \in \mathbb{R}^{F \times F'}$ is a shared and learnable *weight matrix*, and σ is an *activation function* (e.g. logistic/tanh/ReLU)—ignoring biases.
- ▶ The final layer will use the *softmax* function and optimise the *cross-entropy* loss in each training node (usual classification).
- ▶ Simple, but very cheap (and should always be a baseline)!

Augmenting the per-node classifier

- ▶ Many earlier approaches to incorporating graph structure will retain the per-node shared classifier, but incorporate graph structure by either:
 - ▶ *constraining its learnt features* depending on the graph edges;
 - ▶ *augmenting the input layer* with structural node features.
- ▶ I will now briefly cover both of those approaches.

Injecting structure: *semi-supervised embedding*

- ▶ Introduced by Weston *et al.* (ICML 2008), generalising the work of Zhu *et al.* (ICML 2003) and Belkin *et al.* (JMLR 2006) to neural networks.



- ▶ Under the assumption that the edges encode *node similarity*, further constrain the learnt representations of nodes to be close/distant depending on presence of edge!

Semi-supervised embedding loss

- Essentially, the loss function to optimise is augmented with a (dis)similarity constraint, \mathcal{L}_{sim} :

$$\mathcal{L} = \mathcal{L}_0 + \lambda \mathcal{L}_{sim}$$

where \mathcal{L}_0 is the usual supervised learning loss (e.g. cross-entropy), and λ is a hyperparameter.

- One way to define \mathcal{L}_{sim} :

$$\mathcal{L}_{sim} = \sum_i \left(\sum_{j \in \mathcal{N}_i} \|\vec{h}_i - \vec{h}_j\|^2 + \sum_{j \notin \mathcal{N}_i} \max(0, m - \|\vec{h}_i - \vec{h}_j\|^2) \right)$$

where \mathcal{N}_i is the *neighbourhood* of node i , \vec{h}_i is (one of) its hidden layer's outputs, and m is a hyperparameter.

Inserting structure: *DeepWalk*

- ▶ An alternative to augmenting the loss function is first learning some **structural features**, $\vec{\Phi}_i$, for each node i (these will not depend on \vec{f}_i , but on the graph structure)!
- ▶ Then, use $\vec{f}_i \parallel \vec{\Phi}_i$ as the input to the shared classifier (where \parallel is concatenation).
- ▶ Typically, **random walks** are used as the primary input for analysing the structural information of each node.
- ▶ The first method to leverage random walks efficiently is *DeepWalk* by Perozzi *et al.* (KDD 2014)

Overview of DeepWalk

- ▶ Start by random features $\vec{\Phi}_i$ for each node i .
- ▶ Sample a random walk \mathcal{W}_i , starting from node i .
- ▶ For node x at step j , $x = \mathcal{W}_i[j]$, and a node y at step $k \in [j - w, j + w]$, $y = \mathcal{W}_i[k]$, modify $\vec{\Phi}_x$ to maximise $\log \mathbb{P}(y|\vec{\Phi}_x)$ (obtained from a neural network classifier).
- ▶ Inspired by **skip-gram models** in natural language processing: to obtain a good vector representation of a word, its vector should allow us to easily predict the words that *surround* it.

Overview of DeepWalk, *cont'd*

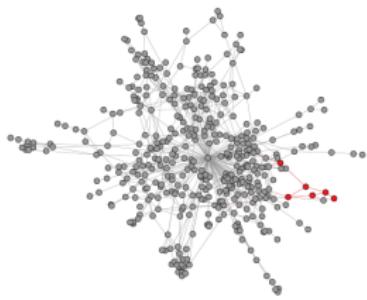
- ▶ Expressing the full $\mathbb{P}(y|\vec{\Phi}_x)$ distribution directly, even for a single layer neural network, where

$$\mathbb{P}(y|\vec{\Phi}_x) = \text{softmax}(\vec{w}_y^T \vec{\Phi}_x) = \frac{\exp(\vec{w}_y^T \vec{\Phi}_x)}{\sum_z \exp(\vec{w}_z^T \vec{\Phi}_x)}$$

is prohibitive for large graphs, as we need to normalise across the entire space of nodes—making most updates *vanish*.

- ▶ To rectify, DeepWalk expresses it as a *hierarchical softmax*—a tree of binary classifiers, each halving the node space.

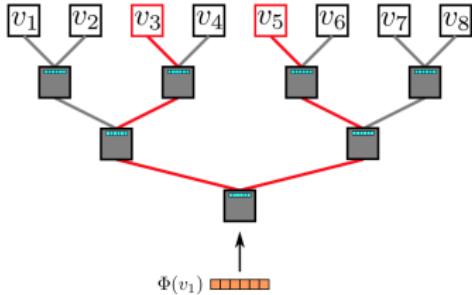
DeepWalk in action



(a) Random walk generation.

$$\mathcal{W}_{v_4} = \begin{bmatrix} 4 \\ 3 \\ 1 \\ 5 \\ 1 \\ \vdots \end{bmatrix} v_j \longrightarrow \Phi^d_j$$

(b) Representation mapping.



(c) Hierarchical Softmax.

Figure 3: Overview of DEEPWALK. We slide a window of length $2w + 1$ over the random walk \mathcal{W}_{v_4} , mapping the central vertex v_1 to its representation $\Phi(v_1)$. Hierarchical Softmax factors out $\Pr(v_3 | \Phi(v_1))$ and $\Pr(v_5 | \Phi(v_1))$ over sequences of probability distributions corresponding to the paths starting at the root and ending at v_3 and v_5 . The representation Φ is updated to maximize the probability of v_1 co-occurring with its context $\{v_3, v_5\}$.

Later improved by *LINE* (Tang *et al.*, WWW 2015) and *node2vec* (Grover & Leskovec, KDD 2016), but main idea stays the same.

Incorporating labels and features: *Planetoid*

- ▶ Methods such as DeepWalk are still favourable when dealing with *fully unsupervised* graph problems, as they don't depend on having any labels or features in the nodes!
- ▶ However, if we have labels/features, **why not use them?**
- ▶ The essence behind **Planetoid** (*Predicting Labels And Neighbours with Embeddings Transductively Or Inductively from Data*), by Yang *et al.* (ICML 2016).

Planetoid's sampling strategy: *Negative sampling*

- ▶ Addresses the issue with $\mathbb{P}(y|\vec{\Phi}_x)$ by employing **negative sampling**; predict instead $\mathbb{P}(\gamma|\vec{\Phi}_x, \vec{w}_y)$, where $\gamma \in \{0, 1\}$.
- ▶ Essentially, use a binary classifier:

$$\mathbb{P}(\gamma|\vec{\Phi}_x, \vec{w}_y) = \sigma \left(\vec{w}_y^T \vec{\Phi}_x \right)$$

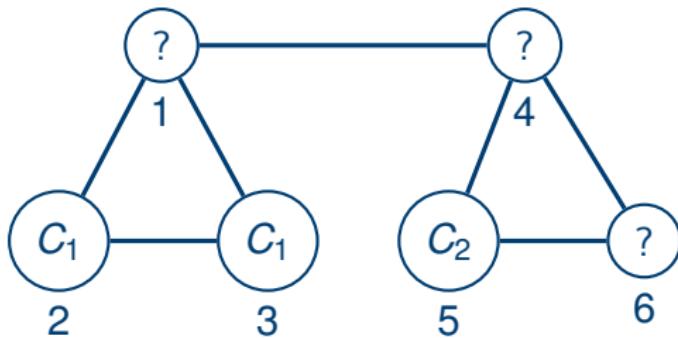
where σ is the logistic sigmoid function. Now each update will focus only on *one* node's weight vector rather than all of them!

- ▶ $\gamma = 1$ implies that nodes x and y are a “positive” pair (more detail in the next slide).

Planetoid's sampling strategy: *Sampling pairs*

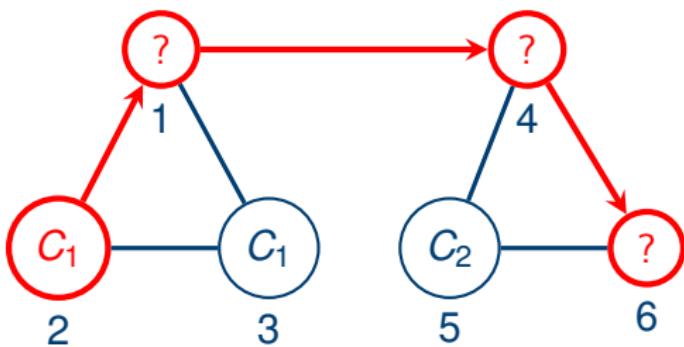
- ▶ Planetoid retains DeepWalk's idea of predicting proximal nodes in random walks.
 - ▶ Sample two nodes a and b that are close enough in a random walk, optimise the classifier to predict $\gamma = 1$.
 - ▶ Sample two nodes a and b uniformly at random, optimise the classifier to predict $\gamma = 0$.
- ▶ It also injects **label information**:
 - ▶ Sample two nodes a and b with same labels ($y_a = y_b$), optimise the classifier to predict $\gamma = 1$.
 - ▶ Sample two nodes a and b with different labels ($y_a \neq y_b$), optimise the classifier to predict $\gamma = 0$.

Planetoid in action



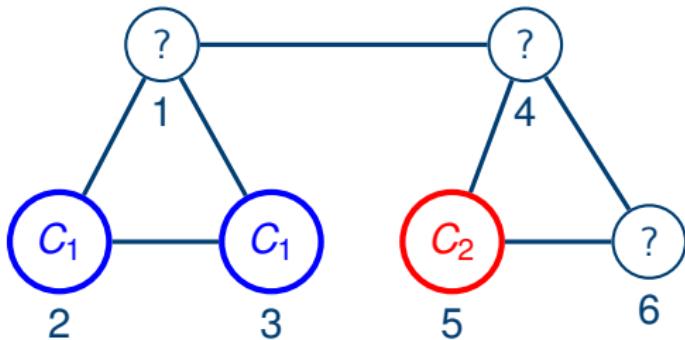
Consider this example graph, with three labelled nodes.
I will now illustrate the two phases of Planetoid.

Planetoid in action: Random walk-based sampling



Sample from a random walk—can take e.g. nodes 1 and 4 with $\gamma = 1$, and nodes 1 and 5 with $\gamma = 0$.

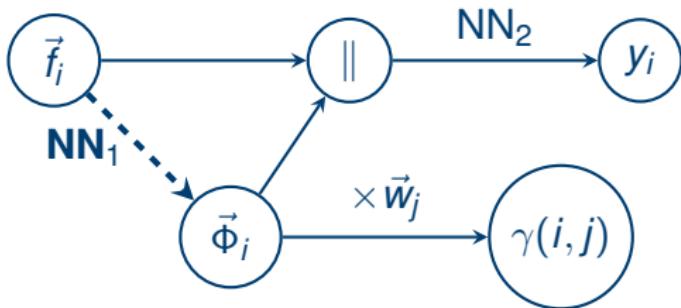
Planetoid in action: Label-based sampling



Sample given the labels—can take e.g. nodes 2 and 3 with $\gamma = 1$,
and nodes 3 and 5 with $\gamma = 0$.

Planetoid's inductive dataflow

- ▶ In an inductive setting, the structural features $\vec{\Phi}_i$ can no longer be independently learned—need to adapt to **unseen** nodes!
- ▶ The inductive version of Planetoid forces $\vec{\Phi}_i$ to directly depend on \vec{f}_i —you guessed it—by employing a neural network. :)



Explicit graph neural network methodologies

- ▶ All methods covered so far have used a shared classifier that classifies each node independently, with graph structure injected only *indirectly*.
- ▶ We will from now restrict our attention solely to methods that *directly* leverage the graph structure when computing intermediate features.
- ▶ **Main idea:** Compute node representations \vec{h}_i based on the initial features \vec{f}_i and the graph structure, and then use \vec{h}_i to classify each node independently (as before).

Graph Neural Networks

- ▶ The first prominent example of such an architecture are **Graph Neural Networks** (GNNs) presented first in Gori *et al.* (IJCNN 2005) and then in Scarselli *et al.* (TNNLS 2009).
- ▶ Start with randomly initialised $\vec{h}_i^{(0)}$, then at each timestep propagate as follows (slightly different than original paper, assuming only undirected edges of one type):

$$\vec{h}_i^{(t)} = \sum_{j \in \mathcal{N}_i} f\left(\vec{h}_j^{(t-1)}\right)$$

where f is a *propagation model*, expressed as a usual neural network linear layer:

$$f(\vec{h}_i) = \mathbf{W}\vec{h}_i + \vec{b}$$

where \mathbf{W} and \vec{b} are learnable weights and biases, respectively.

Graph Neural Networks, *cont'd*

- ▶ As backpropagating through time is expensive, the authors of GNNs further constrain f to be a **contractive map**. This implies that the \vec{h}_i vectors will always converge to a *unique fixed point*!
- ▶ Iterate until convergence (for T steps), then classify using $\vec{h}_i^{(T)}$. Train using the Almeida-Pineda extension of backpropagation (Almeida, 1990; Pineda, 1987).
- ▶ Arguably, too restrictive. Also, impossible to inject problem-specific information into $\vec{h}_i^{(0)}$ (as will always converge to same value regardless of initialisation).

Gated Graph Neural Networks

- ▶ An extension to GNNs, known as **Gated** Graph Neural Networks (GGNNs) by Li *et al.* (ICLR 2016), brought the bleeding-edge deep learning practices to GNNs.
- ▶ Propagate for a **fixed number of steps**, and do not restrict the propagation model to be contractive.
 - ▶ This enables conventional backpropagation.
 - ▶ It also allows us to meaningfully initialise the model!
- ▶ Leverage a more sophisticated propagation model (employing techniques such as *gating*) to surpass GNN performance.

GGNN propagation rule

- ▶ Initialise as $\vec{h}_i^{(0)} = \vec{f}_i \parallel \vec{0}$ (append zeroes for extra capacity).
- ▶ Then propagate as follows (slightly different than original paper, assuming only undirected edges of one type):

$$\vec{a}_i^{(t)} = b_i + \sum_{j \in \mathcal{N}_i} \vec{h}_j^{(t-1)}$$

$$\vec{h}_i^{(t)} = \tanh(\mathbf{W}\vec{a}_i^{(t)})$$

- ▶ Now, extend this to incorporate *gating mechanisms*, to prevent full overwrite of $\vec{h}_i^{(t-1)}$ by $\vec{h}_i^{(t)}$.
 - ▶ Basically, learn (from $\vec{a}_i^{(t)}$ and $\vec{h}_i^{(t-1)}$) how much to overwrite.

Full GGNN propagation rule

- The full propagation model is as follows:

$$\vec{a}_i^{(t)} = b_i + \sum_{j \in \mathcal{N}_i} \vec{h}_j^{(t-1)}$$

$$\vec{r}_i^{(t)} = \sigma \left(\mathbf{W}^r \vec{a}_i^{(t)} + \mathbf{U}^r \vec{h}_i^{(t-1)} \right)$$

$$\vec{z}_i^{(t)} = \sigma \left(\mathbf{W}^z \vec{a}_i^{(t)} + \mathbf{U}^z \vec{h}_i^{(t-1)} \right)$$

$$\vec{\tilde{h}}_i^{(t)} = \tanh \left(\mathbf{W} \vec{a}_i^{(t)} + \mathbf{U} \left(\vec{r}_i^{(t)} \odot \vec{h}_i^{(t-1)} \right) \right)$$

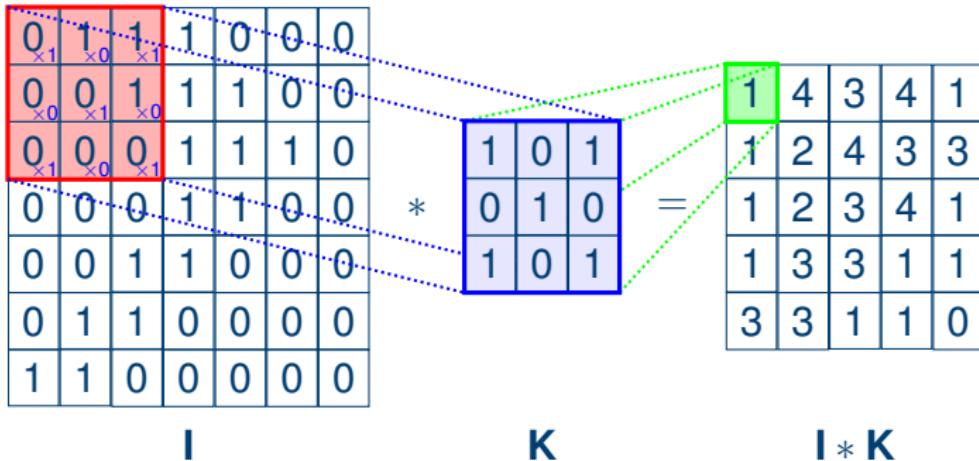
$$\vec{h}_i^{(t)} = (1 - \vec{z}_i^{(t)}) \odot \vec{h}_i^{(t-1)} + \vec{z}_i^{(t)} \odot \vec{\tilde{h}}_i^{(t)}$$

where \odot is elementwise vector multiplication, \vec{r}_i and \vec{z}_i are *reset* and *update* gates, and σ is the logistic sigmoid function.

The silver bullet—a *convolutional* layer

- ▶ GGNNs feature a “time-step” operation which should be very familiar to those of you who have already worked with *recurrent neural networks* (such as LSTMs).
- ▶ These are designed for data that changes *sequentially*; however, our graphs have **static** features!
- ▶ It would be more appropriate if we could somehow generalise the *convolutional operator* (as used in CNNs) to operate on arbitrary graphs!
- ▶ An excellent “common framework” for many of the approaches to be listed now has been presented in “Neural Message Passing for Quantum Chemistry”, by Gilmer *et al.* (ICML 2017).

Convolution on images



Convolution on images

The diagram illustrates the convolution operation $I * K$. It shows three matrices: the input matrix I , the kernel matrix K , and the resulting output matrix $I * K$.

The input matrix I is a 7x7 grid of binary values:

0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	1	1	0	0	0	0
1	1	0	0	0	0	0

The kernel matrix K is a 3x3 grid of binary values:

1	0	1
0	1	0
1	0	1

The resulting output matrix $I * K$ is a 5x5 grid of values:

1	4	3	4	1
1	2	4	3	3
1	2	3	4	1
1	3	3	1	1
3	3	1	1	0

Dotted lines connect the highlighted 3x3 submatrix in I (highlighted in red) to the kernel K (highlighted in blue). The result is highlighted in green in the output matrix $I * K$.

Convolution on images

The diagram illustrates the convolution operation $I * K$. It shows three 7x7 input matrix I , a 3x3 kernel matrix K , and the resulting 7x7 output matrix $I * K$.

Input Matrix I :

0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	1	1	0	0	0	0
1	1	0	0	0	0	0

Kernel Matrix K :

1	0	1
0	1	0
1	0	1

Output Matrix $I * K$:

1	4	3	4	1	0	0
1	2	4	3	3	0	0
1	2	3	4	1	0	0
1	3	3	1	1	0	0
3	3	1	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

Dotted lines indicate the receptive field of each output unit. The result of the convolution is highlighted in green.

Convolution on images

The diagram illustrates the convolution operation $I * K$. It shows three matrices: the input matrix I , the kernel matrix K , and the resulting output matrix $I * K$.

Input Matrix I :

0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	1	1	0	0	0	0
1	1	0	0	0	0	0

Kernel Matrix K :

1	0	1
0	1	0
1	0	1

Output Matrix $I * K$:

1	4	3	4	1
1	2	4	3	3
1	2	3	4	1
1	3	3	1	1
3	3	1	1	0

The diagram shows the convolution process. The input matrix I is overlaid with a red 3x3 box representing the receptive field of the top-left element of the kernel K . The kernel K is shown below it. The result of the convolution is the output matrix $I * K$, where each element is the sum of the products of the corresponding elements in the overlapping regions. The result is highlighted with green dashed boxes.

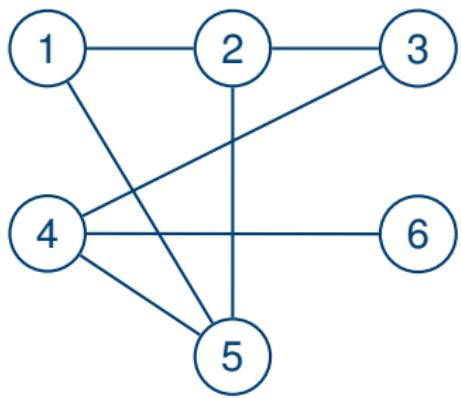
Challenges with graph convolutions

- ▶ Desirable properties for a graph convolutional layer:
 - ▶ **Computational and storage efficiency** ($\sim O(V + E)$);
 - ▶ **Fixed** number of parameters (independent of input size);
 - ▶ **Localisation** (acts on a *local neighbourhood* of a node);
 - ▶ Specifying **different importances** to different neighbours;
 - ▶ Applicability to **inductive problems**.
- ▶ Fortunately, images have a highly rigid and regular connectivity pattern (each pixel “connected” to its eight neighbouring pixels), making such an operator trivial to deploy (as a small kernel matrix which is slided across).
- ▶ Arbitrary graphs are a **much harder** challenge!

Spectral graph convolution

- ▶ A large class of popular approaches attempts to define a convolutional operation by operating on the graph in the **spectral domain**, leveraging the *convolution theorem*.
- ▶ These approaches utilise the **graph Laplacian matrix**, \mathbf{L} , defined as $\mathbf{L} = \mathbf{D} - \mathbf{A}$, where \mathbf{D} is the degree matrix (diagonal matrix with $D_{ii} = \text{deg}(i)$) and \mathbf{A} is the adjacency matrix.
- ▶ Alternately, we may use the **normalised graph Laplacian**, $\tilde{\mathbf{L}} = \mathbf{I} - \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}$.

Graph Laplacian example



$$\mathbf{L} = \begin{bmatrix} 2 & -1 & 0 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & -1 & 0 & -1 & 3 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{bmatrix}$$

Graph Fourier Transform

- ▶ The Laplacian is symmetric and positive semi-definite; we can therefore diagonalise it as $\mathbf{L} = \mathbf{U}\Lambda\mathbf{U}^T$, where Λ is a diagonal matrix of its eigenvalues.
- ▶ This means that multiplying the feature matrix by \mathbf{U}^T allows us to enter the *spectral domain* for the graph! Therein, convolution just amounts to pointwise multiplication.
- ▶ This “Graph Fourier Transform” is the essence of the work of Bruna *et al.* (ICLR 2014).

Graph Fourier Transform, *cont'd*

- ▶ To convolve two signals using the convolution theorem:

$$\text{conv}(\vec{x}, \vec{y}) = \mathbf{U} \left(\mathbf{U}^T \vec{x} \odot \mathbf{U}^T \vec{y} \right)$$

- ▶ Therefore, a *learnable convolutional layer* amounts to:

$$\vec{h}'_i = \mathbf{U} \left(\vec{w} \odot \mathbf{U}^T \mathbf{W} \vec{h}_i \right)$$

where \vec{w} is a learnable vector of weights, and $\mathbf{W} \in \mathbb{R}^{F' \times F}$ is a shared, learnable, feature transformation.

- ▶ Downsides:

- ▶ Computing \mathbf{U} is $O(V^3)$ —*infeasible* for large graphs!
- ▶ One independent weight per node—not fixed!
- ▶ Not localised!

Chebyshev networks

- ▶ These issues have been overcome by *ChebyNets*, the work of Defferrard *et al.* (NIPS 2016).
- ▶ Rather than computing the Fourier transform, use the related family of *Chebyshev polynomials* of order k , T_k :

$$\vec{h}'_i = \sum_{k=0}^K w_k T_k(\mathbf{L}) \mathbf{W} \vec{h}_i$$

- ▶ These polynomials have a recursive definition, highly simplifying the computation:

$$T_0(x) = 1 \quad T_1(x) = x \quad T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$$

Properties of Chebyshev networks

- ▶ Owing to its recursive definition, we can compute the output iteratively as $\sum_{k=0}^K w_k \vec{t}_k$, where:

$$\vec{t}_0 = \mathbf{W}\vec{h}_i \quad \vec{t}_1 = \mathbf{L}\mathbf{W}\vec{h}_i \quad \vec{t}_k = 2\mathbf{L}\vec{t}_{k-1} - \vec{t}_{k-2}$$

where each step constitutes a **sparse** multiplication with \mathbf{L} .

- ▶ The number of parameters is **fixed** (equal to K weights).
- ▶ Note that $T_k(\mathbf{L})$ will be a (weighted) sum of all powers of \mathbf{L} up to \mathbf{L}^k . This means that $T_k(\mathbf{L})_{ij} = 0$ if $dist(i, j) > k$!
⇒ The operator is **K-localised!**

Properties of Chebyshev networks, *cont'd*

- ▶ To avoid issues with exploding or vanishing signals, typically a scaled version of \mathbf{L} is fed into the algorithm:

$$\tilde{\mathbf{L}} = \frac{2\mathbf{L}}{\lambda_{max}} - \mathbf{I}$$

where λ_{max} is the largest eigenvalue of \mathbf{L} .

- ▶ This constrains all eigenvalues to lie in the range $[-1, 1]$, therefore making the norm of all results controllable.
- ▶ Major limitation: *unable to specify **different weights** to **different nodes** in a neighbourhood!* All k -hop neighbours will receive weight $w_k + w_{k+1} + \dots + w_K$.

Limited filters

Going back to the image scenario, under the assumption that each pixel of an image is connected to its immediate four neighbours, this would constrain our 3×3 convolutional kernel to be of the form:

$$\begin{bmatrix} w_2 & w_1 + w_2 & w_2 \\ w_1 + w_2 & w_0 + w_1 + w_2 & w_1 + w_2 \\ w_2 & w_1 + w_2 & w_2 \end{bmatrix}$$

severely limiting the variety of patterns that can be usefully extracted from the image.

GCNs

- ▶ Arguably the most popular approach in recent months has been the **Graph Convolutional Network** (GCN) of Kipf & Welling (ICLR 2017).
- ▶ The authors further simplify the Chebyshev framework, setting $K = 1$ and assuming $\lambda_{max} \approx 2$, allowing them to redefine a single convolutional layer as simply:

$$\vec{h}'_i = \tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2} \mathbf{W} \vec{h}_i$$

which improves computational performance on larger graphs and predictive performance on small training sets.

- ▶ However, the previous issue is *still there...*

Applicability to inductive problems

- ▶ Another *fundamental* constraint of all spectral-based methods is that the learnt filter weights are assuming a particular, fixed, graph Laplacian.
- ▶ This makes them theoretically inadequate for arbitrary **inductive** problems!
- ▶ We have to move on to non-spectral approaches . . .

Molecular fingerprinting networks

- ▶ An early notable approach towards such methods is the work of Duvenaud *et al.* (NIPS 2015).
- ▶ Here, the method adapts to processing with various degrees by learning a *separate* weight matrix \mathbf{H}_d for each node degree d .
- ▶ The authors dealt with an extremely specific domain problem (*molecular fingerprinting*), where node degrees could never exceed five; this **does not scale** to graphs with *very wide degree distributions*.

GraphSAGE

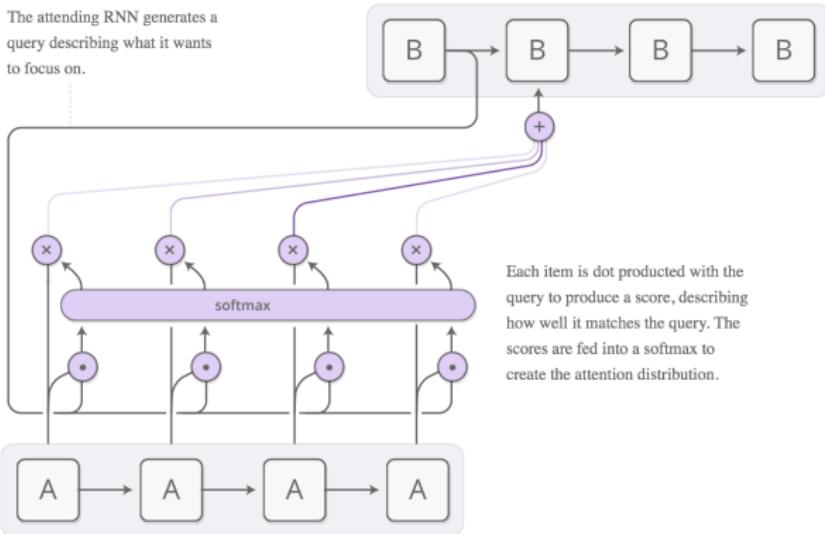
- ▶ Conversely, the recently-published **GraphSAGE** model by Hamilton *et al.* (NIPS 2017) aims to **restrict every degree to be the same** (by sampling a *fixed-size* set of neighbours of every node, during both training and inference).
- ▶ Inherently **drops relevant data**—limiting the set of neighbours visible to the algorithm.
- ▶ Impressive performance was achieved across a variety of inductive graph problems. However, the best results were often achieved with an LSTM-based aggregator, which is unlikely to be optimal.

Attentional mechanisms

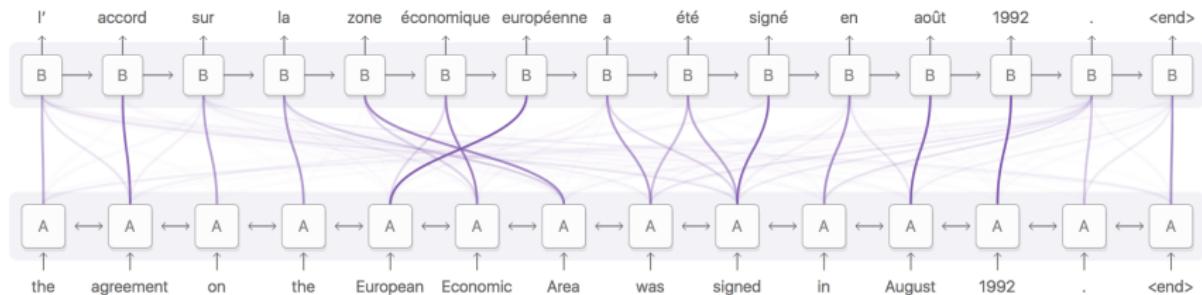
- ▶ One of the latest non-spectral techniques leverages an *attentional mechanism* (originally published by Bahdanau *et al.* (ICLR 2015)), which is now a *de facto* standard for sequential processing tasks.
- ▶ Computes *linear combinations* of the input features to generate the output. The coefficients of these linear combinations are parametrised by a **shared neural network!**
- ▶ Intuitively, allows each component of the output to generate its own combination of the inputs—thus, different outputs *pay different levels of attention* to the respective inputs.

Attention in action: *a potential mechanism*

The attending RNN generates a query describing what it wants to focus on.



Attention in action: *machine translation*



Self-attention

- ▶ A rather exciting development in this direction concerns **self-attention**; a scenario where the input *attends over itself*:

$$\alpha_{ij} = a(\vec{h}_i, \vec{h}_j)$$

$$\vec{h}'_i = \sum_j \text{softmax}_j(\alpha_{ij}) \vec{h}_j$$

where $a(\vec{x}, \vec{y})$ is a neural network (the *attention mechanism*).

- ▶ Critically, this is **parallelisable** across all input positions!
- ▶ Vaswani et al. (NIPS 2017) have successfully demonstrated that this operation is self-sufficient for achieving state-of-the-art on machine translation.

Graph Attention Networks

- ▶ My recent ICLR 2018 publication—in collaboration with the Montréal Institute for Learning Algorithms (MILA)—proposing **Graph Attention Networks** (GATs), leverages exactly the self-attention operator!
- ▶ In its naïve form, the operator would compute attention coefficients *over all pairs of nodes*.
- ▶ To inject the graph structure into the model, we *restrict* the model to only attend over a node's neighbourhood when computing its coefficient!

GAT equations

- ▶ To recap, a single attention head of a GAT model performs the following computation:

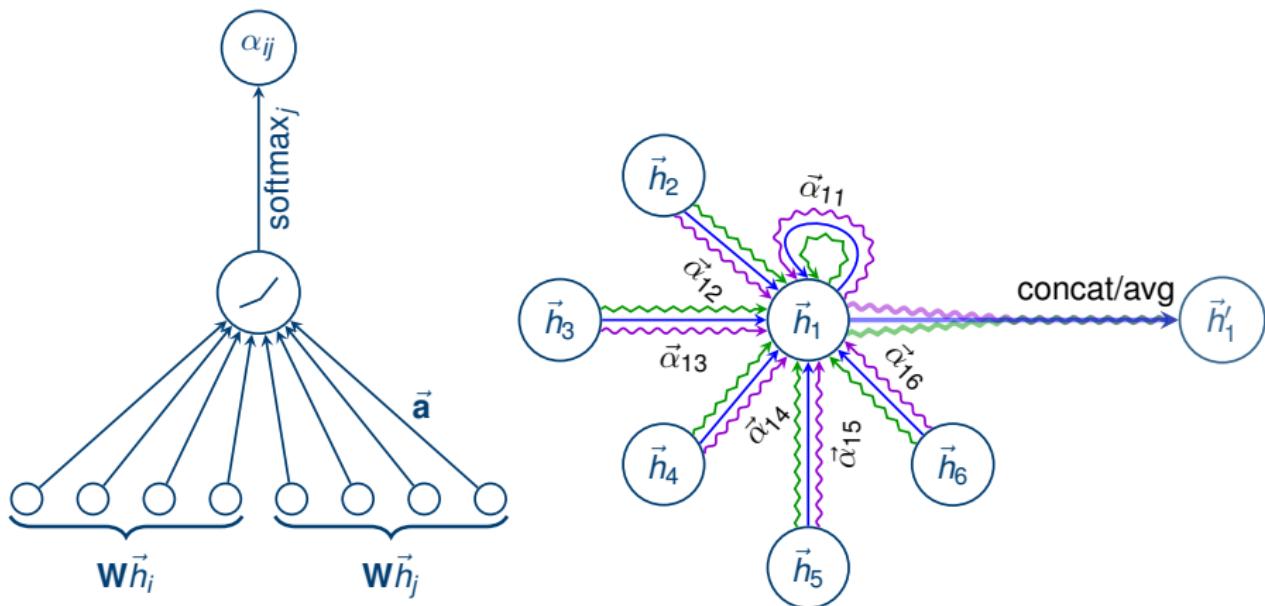
$$e_{ij} = a(\mathbf{W}\vec{h}_i, \mathbf{W}\vec{h}_j)$$

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}_i} \exp(e_{ik})}$$

$$\vec{h}'_i = \sigma \left(\sum_{j \in \mathcal{N}_i} \alpha_{ij} \mathbf{W}\vec{h}_j \right)$$

- ▶ Some further optimisations (like *multi-head attention* and *dropout* on the α_{ij} values) help further *stabilise* and *regularise* the model.

A single GAT step, visualised



GAT analysis

- ▶ **Computationally efficient:** attention computation can be parallelised across all edges of the graph, and aggregation across all nodes!
- ▶ **Storage efficient**—a sparse version does not require storing more than $O(V + E)$ entries anywhere;
- ▶ **Fixed** number of parameters (dependent only on the desirable feature count, not on the node count);
- ▶ Trivially **localised** (as we aggregate only over neighbourhoods);
- ▶ Allows for (implicitly) specifying **different importances** to **different neighbours**.
- ▶ Readily applicable to **inductive problems** (as it is a shared *edge-wise* mechanism)!

GAT performance

- ▶ It seems that we have finally satisfied all of the major requirements for our convolution!
- ▶ How well does it perform?

Datasets under study

Table: Summary of the datasets used in our experiments.

	Cora	<i>Transductive</i>		<i>Inductive</i>
		Citeseer	Pubmed	PPI
# Nodes	2708	3327	19717	56944 (24 graphs)
# Edges	5429	4732	44338	818716
# Features/Node	1433	3703	500	50
# Classes	7	6	3	121 (multilabel)
# Training Nodes	140	120	60	44906 (20 graphs)
# Validation Nodes	500	500	500	6514 (2 graphs)
# Test Nodes	1000	1000	1000	5524 (2 graphs)

Results on Cora/Citeseer/Pubmed

Transductive

Method	Cora	Citeseer	Pubmed
MLP	55.1%	46.5%	71.4%
ManiReg	59.5%	60.1%	70.7%
SemiEmb	59.0%	59.6%	71.7%
LP	68.0%	45.3%	63.0%
DeepWalk	67.2%	43.2%	65.3%
ICA	75.1%	69.1%	73.9%
Planetoid	75.7%	64.7%	77.2%
Chebyshev	81.2%	69.8%	74.4%
GCN	81.5%	70.3%	79.0%
MoNet	$81.7 \pm 0.5\%$	—	$78.8 \pm 0.3\%$
GCN-64*	$81.4 \pm 0.5\%$	$70.9 \pm 0.5\%$	79.0 $\pm 0.3\%$
GAT (ours)	83.0 $\pm 0.7\%$	72.5 $\pm 0.7\%$	79.0 $\pm 0.3\%$

Results on PPI

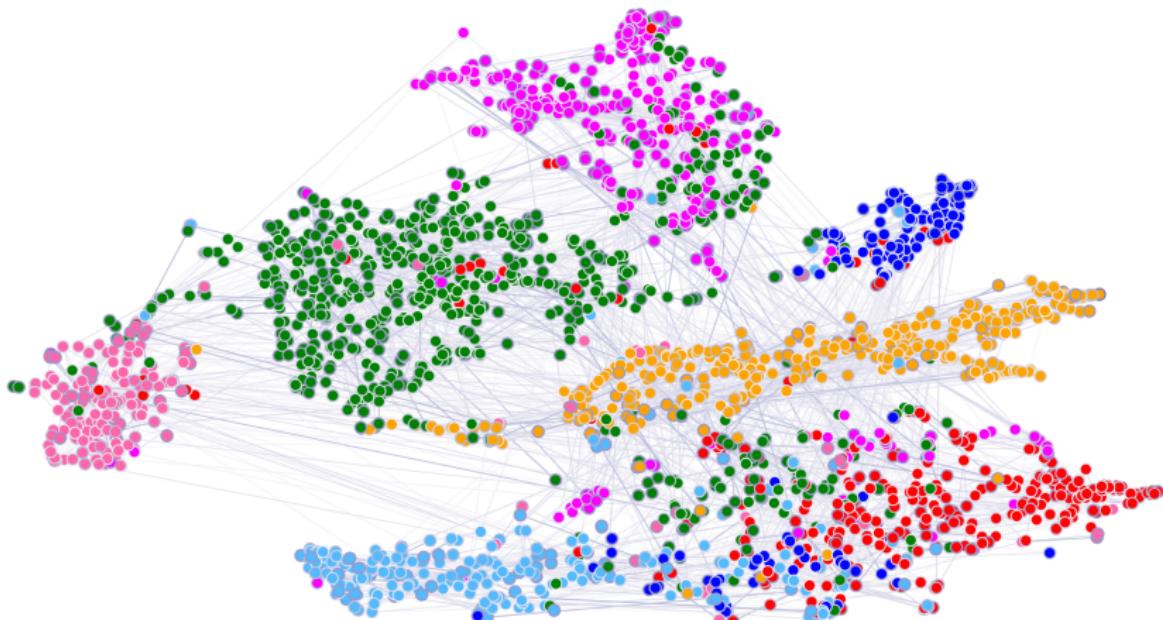
<i>Inductive</i>	
Method	PPI
Random	0.396
MLP	0.422
GraphSAGE-GCN	0.500
GraphSAGE-mean	0.598
GraphSAGE-LSTM	0.612
GraphSAGE-pool	0.600
GraphSAGE*	0.768
Const-GAT (ours)	0.934 ± 0.006
GAT (ours)	0.973 ± 0.002

Here, *Const-GAT* is a GCN-like inductive model.

Applications

- ▶ I will conclude with an overview of a few interesting applications of GCN- and GAT-like models.
- ▶ This list is by no means exhaustive, and represents only what I have been able to find thus far. :)

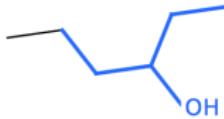
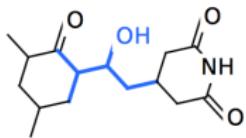
Citation networks



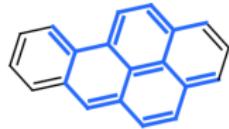
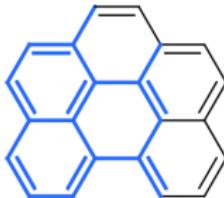
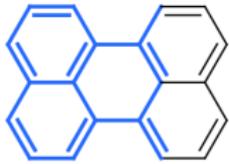
Veličković *et al.* (ICLR 2018)

Molecular fingerprinting

Fragments most activated by pro-solubility feature



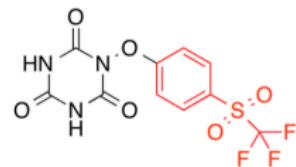
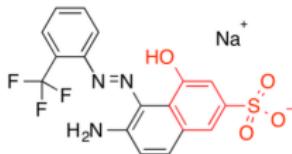
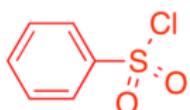
Fragments most activated by anti-solubility feature



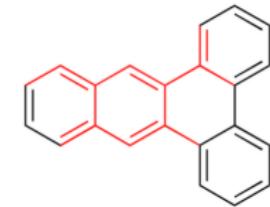
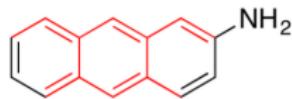
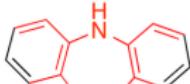
Duvenaud *et al.* (NIPS 2015)

Molecular fingerprinting, *cont'd*

Fragments most activated by toxicity feature on SR-MMP dataset

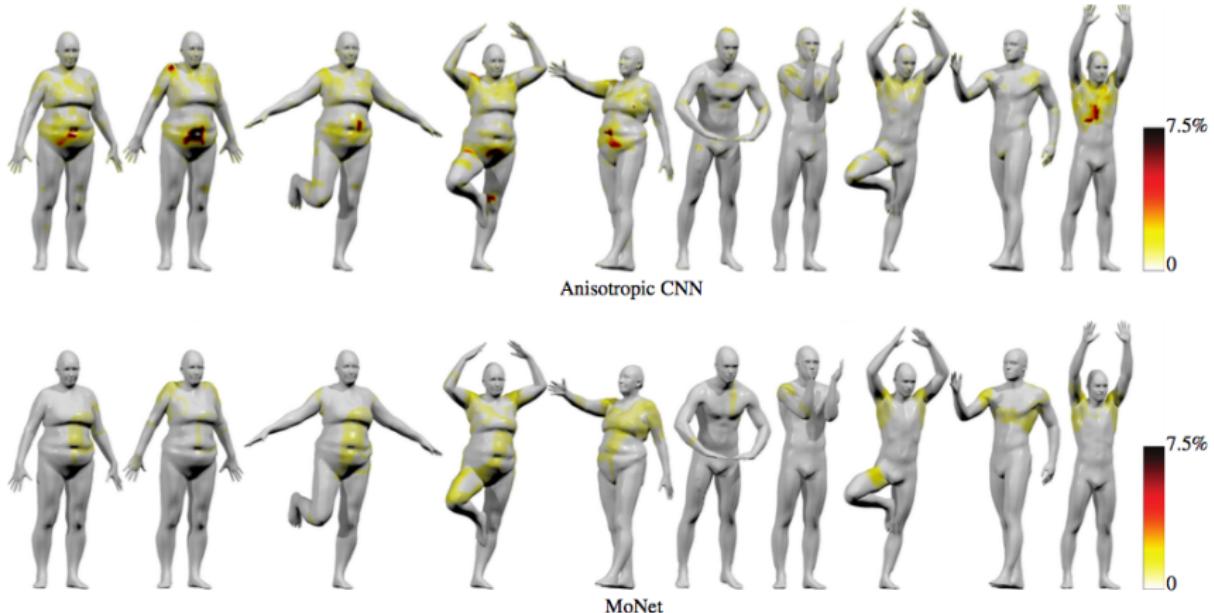


Fragments most activated by toxicity feature on NR-AHR dataset



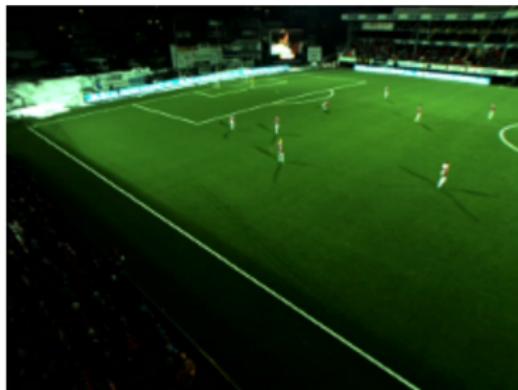
Duvenaud *et al.* (NIPS 2015)

Learning on manifolds



The *MoNet* framework, by Monti *et al.* (CVPR 2017)

Modelling multi-agent interactions



The VAI/N framework, by Hoshen (NIPS 2017)

Cortical mesh segmentation

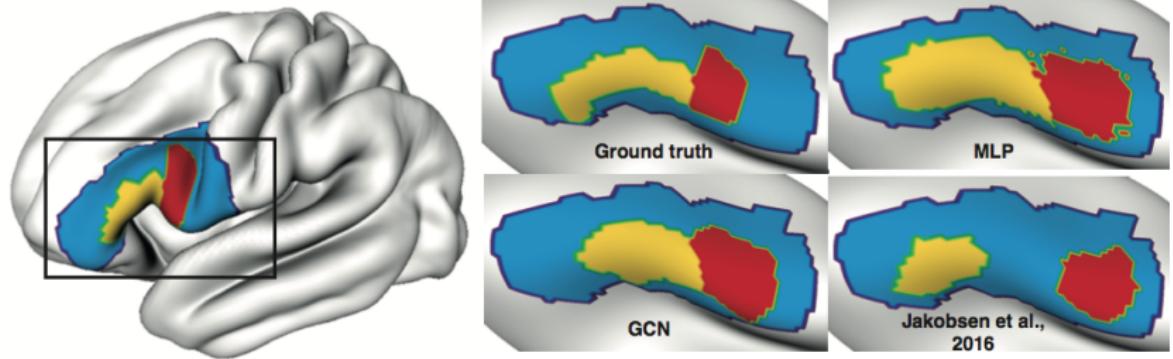


Figure 1: Region segmentation produced by the different models evaluated on the same validation sample.

Cucurull *et al.* (NIPS BigNeuro 2017)

Currently preparing an extended version to submit to MICCAI...

Thank you!

Questions?

petar.velickovic@cst.cam.ac.uk

<http://www.cst.cam.ac.uk/~pv273/>

<https://github.com/PetarV-/GAT>